



THE OPPORTUNITY · SEPTEMBER 2026

The largest carry trade **on earth** is moving.

Thirty years of near-zero Japanese rates built a global funding trade measured in trillions. It is repricing now — and **the information that moves it prints in Japanese while New York sleeps.**

What is in this deck

01 The Opportunity

The trade, the size, and the one actor who moves the price on purpose

03 Primer

Zero-assumption education — every number worked through

05 The Tape

141 weeks of real price data against 28 sourced policy events

07 The Halo

The cross-market confirmation layer — five clusters, five states, one contradiction field

09 The Desk

Nine agents, two humans, and where each sits

11 Build

Ten working days, gate by gate

13 Getting Started

Two entities, the content cut, and what it costs to stand this up

15 Reference

Glossary and specification

02 Game Plan

How a sentence becomes a sized trade, and who does what

04 The Setup

The rate differential, the three forces, the JGB constraint

06 The Edge

The measured window, the escalation ladder, the intervention record

08 Economics

The payoff maths and the three ways this produces return

10 Distribution

90/10, the tiers, and whether the price survives scrutiny

12 Lean Path

Three businesses, what would prove each wrong, and what the answer costs

14 Wiring

Every source, where Sakana AI sits, and what reaches a phone

Where it stands today

Four numbers that frame everything after them.

SPOT

153.61

USD/JPY, 11 Sep close. Down from 164 on 31 July — **6.4% of yen strength in six weeks.**

OFFICIAL FIREPOWER

¥27.1tn

≈ **\$176.6 billion** at 153.61, of yen buying in 2026. The US joined on 3 August, first time since 2011.

CONVEXITY AVAILABLE

3 : 1

Risk \$180k to make \$560k. **Breakeven hit rate 24%.**

TWO REVENUE LINES

\$2.4M

Trading edge plus subscription, on one cost base.

The trade in one paragraph

You borrow yen at 1%, sell it for dollars, and hold something that pays 3.625%. You clear roughly **2 to 2.5% a year in carry** for taking no view on anything.

The position is short yen, unhedged by construction, and **negatively skewed** — many small positive days, rare violent ones. A single 3% move against you erases fourteen months of carry, which is why the unwinds are gap risk rather than trend risk.

Everyone in the trade knows this. What they do not have is a scored, timestamped read on the one actor who can move the price on purpose — Japan's Ministry of Finance — delivered while Tokyo is still trading.

The thesis

One sentence, and the rest of this deck is evidence for it.

Japanese policy language, read in the original against its history and current pricing, improves trading and risk decisions.

A machine that reads it natively and scores it against a measured history of what followed, on **the one actor who moves this price on purpose.**

Not a speed claim. Singapore and Hong Kong trade throughout Tokyo hours, and machine-readable news moves in milliseconds. Section 06 measures when institutions publish relative to when the deepest pools open — which shapes the opportunity. The decision advantage must be earned separately, and can hold even if the speed edge is small.

How a sentence becomes a trade

Four steps, using one real sentence. Most people collapse all four into the word "news" and then wonder why news does not make money.

01 SIGNAL

At 09:47 Tokyo time a Ministry of Finance official said seven words: あらゆる選択肢を排除しない — "we are not ruling out any options." That is all a signal is. A thing that occurred, with a timestamp. **It carries no opinion and is worth nothing on its own.**

02 CONTEXT

Those exact words sit on **rung 6 of an 8-rung ladder**. What followed is on the record: the ministry disclosed **zero intervention** for the window ending 29 July, then **¥15.3993tn — about \$100 billion** in the window opening 30 July. Context is memory, and it is dated. **Wording and spending report on different clocks, and the gap is never assumed.**

03 INTELLIGENCE


Rung 6, third time this year, **with protection still cheap** and positioning crowded on one side. The English wires have not carried it yet. Intelligence is signal plus memory plus **today's price** — the first thing in the chain that could be wrong, and therefore the first that could be valuable.

04 DECISION

Buy a three-month option that pays if the yen strengthens. Risk **0.6% of the book — \$180,000**. If it works it returns roughly **\$740,000**. Above 158.40 the idea was wrong and the position closes. **A decision has a size, a price, and a point at which you admit you were wrong.**

The day, on two clocks

Everything that matters happens while Miami is asleep — and the alert lands in the evening, when the approver is awake.



The alert arrives at about nine in the evening, Miami time. That is the design, not an inconvenience. New York does not see the sentence until half past nine the next morning — twelve hours and forty-three minutes later.

Who does what

Nobody is watching a screen all day waiting for something to happen.

OSAKA — the Observer

Two hours a day, 09:00-11:00 JST. Watches the ministry press pen and the morning wires — the window the measurement says is worth seven hours. Records the words said, who said them, where and at what minute, and how it differed from last time.

The whole job: record what was said, exactly, and never say what you think it means.

THE MACHINES — all night

Twenty-four hours, no handover. They poll, read, classify, score and compose. They never sleep, never get bored on the ninth quiet Tuesday, and never decide this one is probably nothing.

The whole job: turn a sentence into a scored, timestamped, five-line message.

MIAMI — the Approver

Two to five decisions a week, on a phone. Reads five lines. Taps one of three buttons. Nothing reaches the market and nothing reaches a subscriber without that tap.

The whole job: decide, and live with it on the record.

What good and bad look like

A good week

Two alerts. One acted on. Neither wrong in a way that cost anything. The score rose ahead of the move rather than alongside it, and the Japanese read landed before the English wires by a measurable number of minutes.

Quiet weeks are not failures. Most weeks nothing is close to rung 6, the alert budget goes unspent, and the correct output is silence.

A bad week

Nine alerts. That is the failure mode that kills systems like this — not being wrong, but being noisy until the approver stops reading. The budget of two tier-1 alerts a month is enforced in code, not in discipline.

The other bad week: the score moved *after* the price did. That is a lagging indicator in costume, and the day-7 replay is built to catch it before anyone pays for it.

LEAD

minutes

Japanese read timestamp against English wire timestamp.
Measured every time.

NOISE

$\leq 2/\text{mo}$

Tier-1 alerts in a normal regime. Above that, the thresholds are wrong.

HIT RATE

$> 24\%$

Breakeven for the option structure. Wrong three times in four and the year is still flat.

Start from zero

Nothing assumed. The convention that trips up almost everyone comes first.

USD/JPY at 153.61 means one dollar buys 153.61 yen. When that number goes **up**, the yen is getting **weaker**. Yen strength is the rate falling. Every chart in this deck follows that convention.

THE CARRY, WORKED	
Borrow yen at	1.000%
Hold dollars at	3.625%
Gross carry	2.625%
Less funding spread and costs	~0.4%
Cleared, per year, for no view	2.0–2.5%
On \$10m of position that is \$200,000 to \$250,000 a year for holding an opinion about nothing at all. Which is precisely why so much capital ends up in it.	

Why the hedged version earns nothing

If the carry trade were free money, everyone would hedge the currency risk and take it. Covered interest parity prices that away exactly.

FORWARD RATE — $F = S \times (1 + R_{JPY}) / (1 + R_{USD})$

Spot 153.61

× (1 + 0.0100)

÷ (1 + 0.03625)

Forward 149.72

Hedged round trip on ¥10m ¥962 ≈ \$6

Six dollars. The forward market has already priced the entire interest advantage away. The only way to earn the carry is to leave the currency unhedged — which means the carry is not a free lunch, it is **payment for taking gap risk.**

The shape of the risk

Many small gains. Rare violent losses. That is the pattern, not bad luck.

What 2.5% a year buys you

Roughly **0.21% a month**, arriving smoothly. It feels like a yield. It is priced like insurance premium, and it behaves like insurance premium.

A single 3% move against the position erases fourteen months of it.

August 2024, for scale

Spot moved **−7.6% in one week** — the largest weekly move in this dataset. Leveraged accounts closed en masse, the Nasdaq 100 sold off hard, and the VIX printed its third-largest spike on record.

The exit is always narrower than the entrance.

This is why the instrument is an option, not a position. Buying the gap rather than selling it inverts the skew: the most you can lose is the premium, decided before you enter, and being wrong repeatedly becomes survivable rather than fatal.

The convexity, worked

Why a 24% hit rate is enough.

ONE POSITION

Premium paid — 0.6% of a \$30m book **\$180,000**

Value if the thesis pays **\$740,000**

Net gain **\$560,000**

Ratio **3.1 : 1**

BREAKEVEN HIT RATE

$180 \div (560 + 180) = 24.3\%$

Wrong three times in four and the year is still flat. That is the whole reason a signal that is right slightly more often than a coin is worth building a company around — and why the measured backtest on the next section does not have to be impressive to be valuable.

The differential does not close this week

Both central banks are moving in the same direction, which leaves the gap where it was.

IF BOTH HIKE 25BP THIS WEEK		
	BEFORE	AFTER
US policy rate	3.625%	3.875%
Japan policy rate	1.000%	1.250%
Differential	2.625%	2.625%
Identical. The reason to hold the trade is unchanged by the meetings everyone is watching. The move is therefore in the surprise — the vote split, the dissents, the guidance wording — not in the level. That is a point most coverage this week will miss.		

Which is exactly the kind of thing a system reading the statement in the original language is built to catch, and a headline feed is not.

Three forces, and what each one decides

Rate differential

The reason the trade exists. Roughly 2.6% and currently stable. **Decision:** a narrowing differential removes the reason to hold; a stable one removes the urgency to exit and lets positioning build.

Official intervention

The only actor that moves this price deliberately. ¥27.134tn (≈\$176.6bn at 153.61) deployed in 2026. **Decision:** escalation wording is the leading indicator; the spend is confirmation that arrives weeks later.

Positioning and crowding

How many are on the same side and how fast they can leave. **Decision:** crowding does not cause the move — it determines whether the move becomes an unwind.

141 weeks of real price data

ECB reference rates, 2 Jan 2024 to 7 Sep 2026, against 28 sourced policy events. The ECB describes these as informational and discourages transaction use — history and context, not execution data.



What the tape actually did afterwards

The measured result, computed rather than chosen. It is not flattering, and it is kept.

EVENT TYPE	N	AVG +4WK	YEN STRONGER AT +4WK
Intervention	7	-1.54%	3 of 7 (43%)
Joint action	4	+0.61%	1 of 4 (25%)
Japan policy	10	-0.55%	6 of 10 (60%)
US policy	5	+0.25%	2 of 5 (40%)
All events	28	-0.40%	13 of 28 (46%)

Two limits on what this table supports. The samples are small: six of ten gives an approximate 95% interval of **31%-83%**, which does not establish that any subset is meaningfully above even. And **a stronger yen at four weeks is not a profitable option** — payoff depends on strike, entry volatility, elapsed time, exit price and costs. Direction at a fixed horizon is a diagnostic, not a return.

What the table is telling you

Japan policy events read best

6 of 10 followed by a stronger yen at four weeks, on a 95% interval of 31%-83%. Suggestive, not established. Statements and decisions rather than operations — the class this system reads first.

Interventions fade

3 of 7, averaging -1.54% . Buying yen after the ministry has already bought it is late. The operation is the confirmation, not the signal — which is why the escalation wording matters more than the spend.

Joint action is different in kind

+0.61% average, the only positive class. Small n , but co-ordinated operations behave unlike unilateral ones and cannot be faded on the same assumptions.

What this does not say

It measures direction at a fixed horizon. It does not measure the option payoff, which cares about **magnitude and timing** — and that is where a 3:1 structure earns its keep.

The window, measured

The lead this system claims had been asserted from time zones. It has now been computed.

MEDIAN → NEW YORK

10.5 h

Range 9.0 to 12.5 hours. n = 11.

MEDIAN → LONDON

5.5 h

Range 4.0 to 7.5. **London is the real competitor.**

BEST CLASS

7.0 h

Morning ministry remarks, before London opens

WORST CLASS

0.5 h

The 15:30 Governor's press conference

EVENT CLASS

HOURS BEFORE LONDON OPENS

→ NEW YORK

Ministry remarks

09:30 JST · n=4

Rate check


10:00 JST · n=1

BOJ decision

12:00 JST · n=5

Governor's conference

15:30 JST



11.75h

12.00h

9.60h

5.50h

The finding, and the correction it forced

EVENT CLASS	JST	N	→ LONDON	→ NEW YORK	WHAT IS OPEN MEANWHILE
Ministry remarks	09:30	4	7.00h	11.75h	Tokyo only
Rate check	10:00	1	7.00h	12.00h	Tokyo only
BOJ board speech	10:30	1	5.50h	10.50h	Tokyo only
BOJ policy decision	12:00	5	4.60h	9.60h	Tokyo, then Asia
Governor's conference	15:30	—	0.50h	5.50h	Tokyo close, London imminent
All Japanese events	—	11	5.50h med	10.50h med	—
CONTROL — US-ORIGIN EVENTS, MEASURED FORWARD TO THE TOKYO OPEN					
FOMC decisions	14:00 ET	4	6.00h to Tokyo	New York and London both trading	
Jackson Hole	10:00 ET	1	10.00h to Tokyo	New York and London both trading	

This corrected the operating design. The Observer had been scheduled for the Tokyo afternoon, around the press conference — on these numbers the worst two hours of the day to be watching. **The Observer's hours moved to 09:00–11:00 JST**, where the window is seven hours and the ministry speaks.

Why the control matters

US events give a mean **6.60 hours** to the Tokyo open — close enough to the Japanese figure against London that, taken alone, it would suggest the asymmetry is imaginary.

It is not, and the reason is what is open during those hours. A Federal Reserve decision at 14:00 New York time lands while New York and London are both trading: it is priced within minutes, and the six hours until Tokyo opens are six hours of an already-adjusted price. A ministry remark at 09:30 Tokyo time lands when **both deep centres are closed**.

The clock gap is similar. The liquidity behind it is not remotely similar — and that difference is the entire product.

Method, coverage, and what is not yet included

Method

Each event placed at its institution's scheduled publication time in JST, converted to UTC, measured forward to the next 08:00 local open in London and New York — **with daylight saving resolved per event date**, worth a full hour for events between late March and late October.

Coverage: 11 of 28 events. The seven interventions are excluded — an operation has no scheduled time and is the outcome the system predicts, not an input. Two market events and three bilateral meetings have no fixed schedule. Five US events form the control.

What this measures — and what it does not

This is the **structural** component: hours between a Japanese institution publishing and the deeper markets opening. Arithmetic on published schedules, so it is exact.

It does **not** include the transmission component — minutes between a Japanese-language source carrying a statement and an English wire carrying it. That needs paired timestamps from both feeds, and it is the first thing the reader gets instrumented to record. **Every classification thereafter produces one observation of it, free.**

The escalation ladder

Japanese official language escalates by degree. The rungs are separated by a single adjective, and the English wires flatten most of them into the same two words.

WHAT WAS ACTUALLY SAID	HOW IT USUALLY ARRIVES IN ENGLISH	RUNG	WHAT IT SIGNALS
注視している	"watching closely"	1	Routine. Said most weeks. No information.
緊張感を持って注視	"watching closely"	3	A degree added. The ladder has started.
高い緊張感を持って	"watching closely"	4	Another degree. Still the same three English words.
一方的で急激な動き	"one-sided moves"	5	The formal justification clause, laid down for later.
あらゆる選択肢を排除しない	"all options on the table"	6	The rung that has historically preceded action.
断固たる措置をとる	"decisive action"	7	The last thing said before money moves.

Three of those six arrive in English as the same phrase. A model reading a translation cannot recover the difference, because the difference was deleted before it got there.

Why the Japanese path runs first



Axis is ordered, not to scale. The same sentence reaches both paths at the same instant; what differs is how long it takes to become a decision. The system does not need to be right more often than the wires — it needs to arrive before they do.

The halo is real, and it is rare

Forty-five dated yen events. 1,482 market-window observations across fourteen public series. One mechanical rule applied to every one of them. Five events pass.

EVENTS MEASURED

45

22 Sep 1985 to 15 Sep 2026, against 1,482 one-, five- and twenty-day windows.

MEETING THE RULE

5

11.1%. Three of the five are the same week in August 2024.

BOJ DECISIONS

0.47

Mean breadth out of 8, across 19 policy decisions. **Not one has ever reached 3.**

COST TO BUILD IT


\$0–2.4k

A year. Every series is published free by a central bank or an exchange.

Here is what that buys you. Roughly two hundred and forty trading days a year throw off something that looks like a yen event. Four of them are. This layer separates those four from the rest, then says which of the four you are looking at — a Japan-only repricing, a global growth shock, a positioning unwind, or a funding squeeze. **Four different trades, four different holding periods, and the same first thirty minutes on the screen.**


Forty-one years of measurement, on one page

Eight binary one-day flags: yen strength, carry-FX weakness, equity selloff, volatility spike, duration bid, credit widening, oil weakness, crypto deleveraging.



The halo is not made in Tokyo

The family with the largest halo is American data. The family with the most events has the second-smallest halo on the page.



Where the yen actually goes afterwards

Breadth is a snapshot. What the price does over the following month separates the families very cleanly.

FAMILY	N	1 DAY	5 DAYS	20 DAYS	READING
FX intervention	14	+0.25%	+0.80%	+1.27%	The one persistent pattern in the file. Yen strength that builds rather than fades
External catalyst	2	+1.76%	+3.55%	+0.49%	Violent, then largely retraced. n=2 — an illustration, not a finding
Funding stress	1	+2.06%	+0.33%	+0.70%	Gave almost all of it back within a week
BOJ policy	19	+0.04%	-0.19%	+0.09%	Nothing, at any horizon. Nineteen decisions, no measurable drift


The primary signal sits on the only family with a persistent yen path. Reading official language to anticipate *intervention* is where the edge is claimed, and the measurement is consistent with it. **A Bank of Japan decision is context, not a trading event** — nineteen of them say so.

DECISION — Scheduled BOJ decisions are downgraded from alert-generating events to context. They set the funding regime and update the ladder’s baseline; they do not on their own raise the stress score. The alert budget is spent on ministry language and on intervention, where the measured path exists.

Why a yen move shows up in markets with nothing to do with Japan

Four investors borrowed the same cheap currency and bought four unrelated things. What overlaps is the loan.

THE CHAIN — ACTIVE ON SOME OCCASIONS, NOT ALL



THREE CHANNELS — THE JOB IS TO WORK OUT WHICH ONE IS RUNNING TODAY

Carry reversal

Yen is repurchased; the financed assets are sold to repay it.

WATCH breadth · carry basket · yen-call skew

Balance-sheet contagion

Losses cut risk capacity; assets are sold because they are liquid, not weak.

WATCH basis · depth · spreads · vol curves

Common macro shock

One surprise moves currencies, bonds, equities and commodities at once.

WATCH UST/JGB · dollar · VIX · copper · oil

The design rule that makes the rest of this work

Almost every naive scoring system fails here — and it fails in the direction that flatters itself.

A confirmation is worth having only when it brings a **different mechanism**. Equity volatility, equity losses and high-yield spreads cannot be counted as three independent votes when all three are simply the same risk shock measured three ways. Counted naively, one piece of information becomes a high-confidence reading.

The broad dollar belongs in the control set, not the signal set. A dollar shock moves the yen, weakens high-beta currencies and pressures global equities — a pattern that looks exactly like a carry unwind while having nothing to do with Japanese funding.

Three votes, or one?

A day when the VIX jumps, equities fall and high-yield spreads widen looks like **three confirmations**. It is one — a single risk shock, measured three ways, by instruments that are priced off each other. **Scored naively it produces a confident reading built on one observation.**

What a real second vote looks like

Yen-call premium repricing is a *different* mechanism from equity losses: it is somebody paying up for protection in a market that does not have to move for the price to change. **That is new information, not the same information restated.**

Why this matters commercially


A subscriber who discovers that three of your five green lights were the same light will not renew. The independence rule is what makes the confidence number **auditable** — and auditable is what a treasury or a fund can actually buy.

DECISION — The dollar index and the US-Japan rate differential enter the residual calculation as controls on day one, before any cluster is scored. Gold and crypto are admitted to the contradiction field only, and are never permitted to raise confidence.

Five clusters, one vote each

Drawn along mechanism lines, not asset lines. Combined freely within a cluster; added across them.

FIVE MECHANISMS · ONE VOTE EACH



Primary yen signal, plus three independent confirmations, equals a testable high-confidence state —
and the one cluster that refuses to agree stays on the face of the alert, because it is usually the most informative field on it.

The numbers on this diagram are the worked example three slides on, not a live reading. Three clusters confirm, one is neutral, and funding contradicts — precisely the combination that says *broad unwind, but technical rather than systemic*. The contradiction stays on the face of the alert because it is usually the most informative field on it.

What each cluster is for

CLUSTER	CORE MEASURES	QUESTION IT ANSWERS	ROLE
1 · FX breadth	USD/JPY, AUD/JPY, NZD/JPY, MXN/JPY, ZAR/JPY, liquid proxies for BRL and TRY	Is this a broad carry reversal, or one currency with its own problem?	Required anchor
2 · Vol and options	1-week and 1-month JPY implied vol, normalised JPY-call premium, VIX, term structures	Is protection repricing before spot contagion arrives?	Early amplifier
3 · Rates and macro	US-Japan 2-year and OIS differential, JGB front end, UST 2s and 10s, dollar index, surprises	What caused the move?	Regime classifier
4 · Funding and credit	Cross-currency basis, repo proxies, high-yield spreads, credit default swaps, bid-ask and depth	Is balance-sheet stress spreading beyond price?	Systemic gate
5 · Risk assets	Nikkei and TOPIX futures, global high beta, copper and oil, commodity currencies, BTC and ETH	Where is the halo actually landing?	Opportunity map

Every instrument keeps its own clock. Regulatory positioning data is published days late — excellent for judging vulnerability, useless for confirming anything live. A score assembled from live FX, delayed credit and week-old positioning, presented as one number with one timestamp, is four different moments pretending to be one.

Where the halo is strong, and where it is not

Grading each market honestly is what stops the layer confirming whatever it already believes.

MARKET OR MEASURE	EXPECTED RESPONSE IN AN UNWIND	EVIDENCE	ROLE IN THE WATCH
JPY-funded FX basket	Yen rises; high-yield and high-beta currencies weaken	Strong · direct	Primary breadth
JPY implied vol and skew	Vol rises; yen-call protection reprices	Strong · direct	Early confirmation
Japanese and global equities	High-beta and leveraged exposures sell first	Strong	Risk transmission
VIX and equity vol	Spot vol rises; curves may flatten or invert	Strong	Leverage pressure
Government bonds	Yields fall in a growth scare — and rise under a BOJ shock	Strong · conditional	Catalyst classifier
Credit spreads	May widen modestly, and generally lag equities	Moderate	Systemic check
Commodity currencies	AUD, NZD, BRL, MXN and ZAR weaken against the yen	Strong historically	Carry-macro bridge
Oil and copper	Often fall during a growth or dollar shock	Conditional	Macro control
Gold	Can rise as a haven or fall to raise cash — both are common	Ambiguous	Contradiction only
Crypto	Can reflect leverage reduction; trades when nothing else does	Not yen-specific	Secondary sensor
Basis and funding	Widens as dealer risk capacity contracts	Strong mechanism	Systemic confirmation

Two days that looked identical

On 11 July 2024 the Ministry bought yen. On 5 August 2024 the yen was the focal point of a global leveraged unwind. Both produced almost exactly the same move.

11 JULY 2024 · MOF YEN BUYING

+1.95%

One-day yen appreciation. Breadth **4 of 4** carry crosses. Three of eight mechanical flags.

5 AUGUST 2024 · LEVERAGED UNWIND

+2.06%

One-day yen appreciation. Breadth **4 of 4** carry crosses. Five of eight mechanical flags.

THE CARRY BASKET AGAINST THE YEN — BOTH DAYS, FROM PUBLISHED DAILY RATES

$$r_i/JPY = (1 + r_i/USD) / (1 + rJPY/USD) - 1$$

	11 JUL 2024	5 AUG 2024
MXN/JPY	-1.903%	-4.319%
AUD/JPY	-1.489%	-2.620%
NZD/JPY	-1.427%	-2.709%
BRL/JPY	-2.337%	-2.254%
Ct equal-weight	-1.789%	-2.976%
Bt breadth vs JPY	1.00	1.00
breadth measured		
against the DOLLAR	0.25	1.00

Look at the last line. On 11 July three of the four carry currencies *rose* against the dollar. A screen measuring weakness against the dollar would have called the day narrow. Measured against the yen — the funding leg, and the only reference that describes the trade — **all four fell, on both days.**

The field that decided it was the slow one

Every fast indicator agreed. On the noisier measures August looked more serious. The lagging indicator was the one that was right.


ONE-DAY READING	11 JUL 2024 · M09	5 AUG 2024 · M13	WHAT SEPARATES THEM
Yen, trade-weighted	+1.95%	+2.06%	Indistinguishable
Carry basket vs JPY	-1.79%	-2.98%	Both broad; August wider
Breadth vs JPY	1.00	1.00	Identical. Breadth does not discriminate.
S&P 500	-0.88%	-3.00%	August far more violent
Nikkei 225	+0.94%	-12.40%	Japan equities sold only in August
VIX	+0.07 pts	+15.18 pts	August is the volatility event
THE FIELD THAT ACTUALLY SEPARATES THEM — CREDIT, OVER THE FOLLOWING MONTH			
High-yield spread, day 1	+1 bp	+21 bp	August widened on the day
High-yield spread, day 20	+40 bp	-59 bp	Opposite directions
Yen, 20 days later	+8.85%	+0.70%	The July move ran. The August move gave it back

Same first day. Twelve and a half times the follow-through. Credit kept widening after the July intervention and the yen ran another seven percent; credit recompressed after the August unwind and the yen was back where it started within three weeks. **This is the state-03 versus state-05 distinction, and it cost nothing to observe.** Across the twenty-two events with credit coverage, the eight where credit was wider at twenty days averaged +2.26% of yen appreciation against +0.60% for the fourteen where it was tighter — roughly five independent episodes, so a direction to test rather than a result to trade.

Five states, and the transitions between them

The useful output is not how red the screen is. It is where in the sequence the market currently sits.

WATCH THE TRANSITION, NOT THE LEVEL



Two of these five states are reached without any confirmation at all. Counting red screens cannot distinguish state 03 from state 05 — they can look identical on a dashboard, and they call for opposite trades.

State 03 and state 05 are the pair that matters commercially. Both show a broad, confirmed yen move. In the first, forced sellers are still selling. In the second they are largely finished, and the research question has flipped from *how far does this run* to *what retraces first*. The distinction is made by the **slope** of breadth and stress, not their level — which is why the layer stores a history rather than a snapshot.

What each transition requires

State 03 and state 05 look identical on a dashboard and call for opposite trades.

TRANSITION	REQUIRED OBSERVATION	THE FALSE POSITIVE IT IS BUILT TO REJECT
Vulnerable → Drift	Cross-sectional yen strength exceeds what the rate differential predicts	One weak target currency dragging the average
Drift → Confirmed	Options and a separate risk cluster agree independently	Scheduled-data volatility without breadth behind it
Confirmed → Systemic	Funding, credit or market depth deteriorates	An equity drawdown while financing stays entirely orderly
Confirmed → Exhausting	Breadth narrows; skew and basis improve; credit stops worsening	A brief pause inside renewed deleveraging

DECISION — Every cluster carries its own effective age on the face of the alert, and a cluster older than its stated tolerance is scored zero rather than carried forward. Positioning is admitted to the Vulnerable state only and is barred from the confirmation count.

The same yen move, six different meanings

This is the table that turns a signal into a trade selection. The yen rising is the constant across every row.

OBSERVED COMBINATION	MOST LIKELY REGIME	WHAT FOLLOWS FROM IT
Yen up · JGB yields up · US assets stable	Bank of Japan repricing	Japan-centred expressions; the global halo is weak and should not be assumed
Yen up · UST yields down · VIX up · equities and copper down	Growth scare with carry amplification	Broad risk reduction; duration may hedge the position
Yen up · UST yields up · equities down · credit wider	Inflation, policy or liquidity shock	Duration will not hedge this one. Funding matters more than direction
Yen up · carry FX down · VIX up · credit calm	Technical positioning unwind	Potentially fast, and reversible — watch for exhaustion rather than extension
Yen up · gold, crypto and equities down · basis stressed	Collateral liquidation	Cash and market depth dominate every normal correlation
Yen up · commodities rally · volatility subdued	Supply shock or Japan-specific event	Do not label this a global unwind — it is the classic false alarm

Correlation is the symptom. The causal reading comes from three things a correlation matrix throws away: the *order* in which markets moved, the markets that *refused* to confirm, and which balance-sheet channel became stressed.

The intervention record, and the direction nobody checks

The Ministry publishes every operation since April 1991 — date, amount and direction. There are 386 of them.

OPERATIONS ON RECORD

386

Apr 1991 to Jun 2026 · ¥122.1tn, about \$795bn at 153.61

YEN SELLING

338

87.6% of days. ¥80.9tn, pushing the yen down

YEN BUYING

42

10.9% of days. ¥41.1tn — the regime this desk models

YEARS IT HAS HAPPENED

7

1991, 1992, 1997, 1998, 2022, 2024, 2026

Nearly nine-tenths of Japan's documented intervention history runs opposite to this product's thesis. For most of thirty-five years the Ministry has been a seller of yen: ¥20.4tn across ninety-one days in 2003 alone. A model trained on “the Ministry intervened” without a signed direction field learns the wrong sign roughly eighty-eight percent of the time — from the largest and most confident part of the sample.

DECISION — Direction is a required field on every intervention record, and the escalation ladder is defined only over yen-buying operations. The 2010 and 2011 yen-weakening episodes stay in the archive as controls. No event enters a backtest without a sign.

Seven years in thirty-five

Modern yen buying is concentrated, enormous and rare — which is what makes an escalation ladder worth reading at all.

YEN-BUYING EPISODE	DAYS	TOTAL	USD AT 153.61	WHAT IT ESTABLISHES
1991	3	¥0.056tn	\$0.4bn	Token by modern standards
1992	23	¥0.717tn	\$4.7bn	Many small days — the opposite shape to the modern pattern
1997	3	¥1.059tn	\$6.9bn	Asian crisis period
1998	3	¥3.047tn	\$19.8bn	The deleveraging episode the literature still cites
A TWENTY-FOUR YEAR GAP — THEN THE MODERN REGIME				
2022	3	¥9.188tn	\$59.8bn	First yen buying since 1998
2024	4	¥15.323tn	\$99.8bn	29 Apr, 1 May, 11 Jul, 12 Jul
2026 · Apr-May	3	¥11.735tn	\$76.4bn	Daily detail published 7 Aug 2026
2026 · Jul-Aug	—	¥15.399tn	\$100.2bn	Aggregate only. Daily attribution unresolved at the cutoff
2026 total	3 + window	¥27.134tn	\$176.6bn	The largest yen-buying year on record, by a factor of nearly two

Disclosure runs on a different clock from the operation. Daily amounts became knowable 93 to 120 days after execution; the monthly aggregate, 29 days. A backtest using the amount as a feature on the day is reading four months into the future. **DECISION** — every record carries two dates, and the leakage test fails the build if a field is read before its publication timestamp.

Who transmits it — the observability map

An entity belongs here only if it can trigger the event, transmit the balance-sheet adjustment, reveal the positioning, absorb the flow or carry a material earnings consequence.

CLASS	ENTITY OR GROUP	WHAT IS OBSERVABLE	ROLE
OFFICIAL — CAN CAUSE THE EVENT			
Government	Japan Ministry of Finance	Official buying and selling, aggregate and daily amounts, rhetoric on excessive moves	Primary trigger
Central bank	Bank of Japan	Rate decisions, purchase schedule, dissent, financial-conditions language	Funding cost
Central bank	Federal Reserve	Payroll and inflation reaction, two-year yield, swap-line usage	External catalyst
Government group	G7 finance ministries	Joint statements and synchronised dealing	Reversal capacity
INSTITUTIONAL — TRANSMITS OR REVEALS IT			
Institution	Hedge funds, global macro and CTAs	Regulatory yen positioning, carry-index beta, trend flips, prime-broker commentary	Unwind amplifier
Institution	Dealers and global banks	Yen-dollar basis, bid-ask, implied volatility, dealer gamma	Liquidity condition
Institution	Japanese life insurers	Hedge-cost changes, foreign-bond flow, duration reallocation	Slow flow
Institution	Japanese retail FX accounts	Margin balances, leverage rules, Tokyo Financial Exchange positions	Damps, then amplifies
CORPORATE — ABSORBS OR REVEALS THE CONSEQUENCE			
Company basket	Major exporters	FX assumptions, hedge ratios, export mix, guidance	Earnings channel
Company basket	Megabanks	NIM guidance, securities duration, unrealised JGB losses	Orderly vs disorderly
Company basket	Airlines · energy · utilities	Fuel and FX hedge ratios, fuel-adjustment lag, LNG benchmarks	Importer pressure
Company basket	High-beta technology	Nasdaq beta, credit spreads, collateral value	Stress amplifier, not a yen proxy

A map of what is observable — not an accusation that anyone named here caused anything. Technology baskets are crowding barometers, never yen proxies.

What arrives on the phone once the layer is on

The 5 August 2024 reading, reconstructed from published daily series. Four clusters confirmed and the alert would have fired.

TELEGRAM — TIER-1 ALERT, HALO LAYER ENGAGED

▲ **STRESS** 68 → 79 05 Aug 2024 09:40 JST · replay

STATE **CONFIRMED** — broad carry reversal

4 of 5 clusters · 3 of 8 mechanical flags at send

EVIDENCE Breadth 1.00 — all four carry crosses down against
JPY. Basket -2.98%. Yen +2.06%.

CONFIRMS FX breadth · options · risk assets · credit
Nikkei -12.4% · VIX +15.2 · HY +21bp

CONTRA **UST 10y only -2bp. Duration is not confirming
a growth scare.**

CATALYST US labour shock (02 Aug) meeting BOJ repricing
(31 Jul) into crowded positioning.

GATE If HY keeps widening → SYSTEMIC, hold.
If HY recompresses → EXHAUSTING, take profit.

QUALITY All live. BRL/JPY synthetic. Positioning excluded
— 4 days stale.

[**APPROVE**] [**HALVE**] [**DECLINE**]

The field that decided the outcome is GATE — and it did not resolve for three weeks: high-yield spreads recompressed 59 basis points, the state moved to *Exhausting*, and the yen gave back almost all of a two percent move. An alert that only ever presents its supporting evidence is a marketing document; **an alert that names the condition it is waiting on is a research product.**

What it costs, and the order it gets built in

The cheapest phases carry most of the information — because the informative inputs are prices public institutions publish, and the expensive ones are what dealers sell.

PHASE	INPUTS	WHERE THEY COME FROM	ADDED PER YEAR	WHY IN THIS ORDER
1 · Breadth and residual	Ten liquid JPY crosses, US-Japan 2y and OIS, dollar index, event calendar	The intraday FX feed already budgeted, plus public central-bank series	\$0–600	Highest information per dollar in the whole document
2 · Vol and options	JPY implied vol and skew, VIX, Nikkei and TOPIX futures	Exchange settlement data and the public volatility index	\$0–1,200	Reacts before the slower credit data does
3 · Macro classifier	UST and JGB curves, data surprises, copper and oil, commodity FX	Public central-bank and ministry series, plus the same FX feed	\$0–600	Prevents the most expensive error — mislabelling the regime
Phases 1-3 — the recommended build	Everything above, on data already paid for or publicly available		\$0–2,400	Roughly a rounding error on Tier 0
GATED — NOT BUILT UNTIL PHASES 1-3 DEMONSTRATE OUT-OF-SAMPLE LIFT				
4 · Funding and depth	Cross-currency basis, high-yield spreads, credit default swaps, order-book depth	Public spread indices are free; live basis, swaps and depth are dealer products	\$15k–60k	The first genuine step change — and the first needing a data contract
5 · Continuous-market sensor	Crypto spot, perpetual funding, liquidations, weekend proxies	Public exchange endpoints	\$0–300	Cheap, but worth nothing until the primary framework is proven

Phases 1 to 3 are built. 1,482 market-window observations across fourteen series — every carry cross, both equity indices, the Nikkei, the volatility index, the Treasury curve, high-yield spreads, oil, bitcoin and the broad dollar — **all of it from public central-bank and exchange series, at no cost. Tier 0 runs at \$3,950–16,670 a year.** Phases 1 to 3 add **\$0–2,400** — a few hundred dollars a month at the top of the range, nothing at the bottom. **The layer that makes the alert defensible is not the part of this business that costs money.**

How the layer earns its place

The benchmark is not zero. It is the primary yen signal running on its own.

TEST	METHOD	WHAT COUNTS AS PASSING
Lead-lag	One- and five-minute event windows at 5, 15, 30 and 60 minutes	Stable ordering across held-out events
Incremental prediction	Primary-only model against primary plus clusters	Out-of-sample lift in precision or in economic loss
Regime accuracy	Label each episode: BOJ, growth, inflation, liquidity, false alarm	Fewer costly regime confusions than the primary alone
False positives	Matched quiet periods and non-yen shocks	Specificity holds outside the famous episodes
Causality controls	Rates, dollar, China, oil supply and scheduled news held as controls	The carry factor survives every plausible control
Execution realism	Bid-ask, depth, slippage, latency and partial fills applied	The net result is still positive at target size
Capacity	Impact curve by order size and market session	A defined capacity ceiling, and a measured decay
Robustness	Walk-forward estimates, thresholds frozen in advance	No reliance on full-sample normalisation

DECISION — No confidence state is published without its contradiction vector, its per-cluster timestamps and its nearest historical analogue attached. The contradiction field is a required field in the alert schema, and an alert that cannot populate it does not send.

What the evidence supports, and what it does not

CLAIM	WHERE THE EVIDENCE STANDS
Broad FX and equity halo	Well supported — historically, and directly observed in August 2024
Funding and margin amplification	Supported as a mechanism — intensity varies between episodes
Commodity response	Conditional on growth, the dollar, China and supply shocks
Gold response	Ambiguous — haven demand and forced selling pull opposite ways
Crypto as a yen-specific lead	Not established — retained as a risk sensor only
Profitable tradability	Not established by the literature — requires desk testing after costs

The evidence base

BIS Bulletin 90 — market turbulence and the carry trade unwind of August 2024.
BIS Quarterly Review, Sep 2024 — “Carry off, carry on”: the cross-market chronology.
NBER WP 14473 — carry trades and currency crashes; negative skew and funding liquidity.
NY Fed, Q4 1998 — speculative deleveraging as the dollar fell sharply against the yen.
BIS QR, Mar 2007 — AUD, NZD, BRL, TRY and GBP as yen-funded targets.
Fed FEDS Notes, Oct 2023 — separating monetary from growth and risk shocks.
BIS WP 592 — the dollar, bank leverage and covered interest parity deviations.
BIS WP 1000 — dollar beta and stock returns.

What the layer promises is not certainty. It makes uncertainty structured, timestamped and economically interpretable — which is what a sophisticated trader actually buys. Ryūzan produces the observations and the metrics; **ØDN Ventures determines how the feed is packaged, licensed and sold.** The halo layer changes what is measured, not the line between an observation and a recommendation.

Three distinct ways this produces return

The trade itself

Convex positions taken ahead of official action. Risk is capped at premium by construction, so a 24% hit rate breaks even and anything above it compounds.

The subscription

The same signal, published on a stated delay, to an audience that already exists. **Zero marketing cost** — which is why the economics look nothing like a normal subscription build.

The record

A timestamped, immutable public history of every signal raised and what followed. It is the asset that makes the second line defensible and the first line auditable.

These are not three bets on the same outcome. The subscription pays whether or not any individual position works, provided the record is honest. That is what makes the cost base survivable while the trading edge is still being proven.

Hit-rate sensitivity


What the option structure does at each level of being right.

PER 100 POSITIONS · \$180K PREMIUM · \$740K PAYOFF				
Hit rate	Wins	Gross	Premium	Net
15%	15	\$11.1m	\$18.0m	-\$6.9m
24%	24	\$17.8m	\$18.0m	-\$0.2m
30%	30	\$22.2m	\$18.0m	+\$4.2m
40%	40	\$29.6m	\$18.0m	+\$11.6m
46%	46	\$34.0m	\$18.0m	+\$16.0m

The measured direction rate across all 28 events was **46%**, and **60%** on the Japan-policy subset. Those are the rows that matter — and note how far above breakeven both sit. **The structure is doing more work than the signal, and it is meant to.**

Nine agents. Two humans.

Three of the nine are conventional code, not models — thresholds, arithmetic and risk limits cannot be allowed to improvise.




— conventional code — cannot improvise — language model — Japanese-native model — human

Nothing here forms a view. The Watcher does not interpret, the readers classify but do not conclude, the scorer is arithmetic. The Japanese path runs first and the English second, so the interval between the two timestamps is the measured lead.

The decision path

Nothing reaches the market without a person.



The risk gate sits before the human, not after. Anything outside hard limits is killed without ever appearing on a phone, so the approver only sees trades that were already legal to do. The approver is the single point at which a view becomes an action — and the only place in the chain where the answer is not computed.

The two human inputs, and the running cost

THE APPROVER

Five lines on a phone: what changed, the case, the counter-case, the trade, the invalidation. **Approve, decline, or halve.** Two to five decisions a week, covering both the position and its publication.


THE OBSERVER

Osaka, 09:00–11:00 JST — the hours the measurement identifies as the widest window. Records what was said on a fixed four-field form. **There is no recommendation field.**

LINE	DETAIL	ANNUAL
People	Observer part-time plus the principal's own approval time	\$60–175k
Terminal & data	One seat plus programmatic access	\$48–120k
Model layer	Japanese ingestion and the analytical stack	\$30–84k
Platform & tooling	Application platform, store, paging, monitoring, domain	\$3–9k
Steady state	Serving both revenue lines, observer hired	\$140–390k
Before that	Phases 0–2, no observer, no terminal, free data	\$4–12k/yr

Sakana reads it. A person confirms it.

Machine interpretation is the accelerant. It is not the authority.



The signal does not deploy on a model's reading alone. Sakana AI classifies the Japanese statement in seconds and the score recomputes immediately — but an escalation-grade reading goes to a **human Japanese verifier** before anything leaves the building. Ten to thirty items a month need that, paid per item. The machine buys the minutes; the person carries the claim.

90 / 10

Roughly one follower in ten will genuinely act on a trading idea. On an audience of ninety thousand, that tenth is the business.

THE FOLLOWING

90,000

Reach. Real credibility, weak as a planning number.

THE ENGAGED TENTH

9,000

Read closely, act, and would pay. **Nine tenths of the audience's economic value sits here.**

PAYING

1,800

20% of the tenth — ordinary inside an engaged group. Two per cent of total, reached bottom-up.

Plan, price and write for the nine thousand. A product aimed at ninety thousand has to be broad and shallow to reach them. A product aimed at nine thousand people who actually trade can be **narrow, expensive and detailed** — the only kind this machine produces.

What 1,800 subscribers actually looks like

The headline number is only useful if it decomposes. It does.

TIER	MONTHLY	ANNUAL	SUBSCRIBERS	SHARE	REVENUE
Watch — daily read, stress score, same-day translation	\$39	\$468	1,186	66%	\$555,000
Signal — the alerts as they fire, with invalidation	\$149	\$1,788	571	32%	\$1,021,000
Desk — raw observation forms, programmatic score, capped seats	\$999	\$11,988	54	3%	\$644,000
Total	—	—	1,811	100%	\$2,223,348

Average revenue per subscriber is **\$1,227.69 a year — about \$102 a month**. High for a newsletter, ordinary for a professional tool, which is the correct place for this to sit. **Fifty-four institutional seats carry 29% of revenue**, and they are the least likely to leave.

Does the price survive contact with what else exists?

WHAT	WHO IT SERVES	INDICATIVE PRICE
General equity research portals	Retail investors	≈\$300/yr
WATCH SITS HERE — \$468/YR		
Macro research letters	Semi-professional macro readers	\$500–1,000/yr
Active-trader terminals and scanners	Serious retail traders	\$1,200–2,700/yr
SIGNAL SITS HERE — \$1,788/YR		
Independent institutional macro research	Funds and bank desks	\$25,000–60,000/yr
LSEG Workspace	Professional desks	≈\$22,000/seat/yr
Bloomberg Terminal	Professional desks	≈\$30,000+/seat/yr
DESK SITS HERE — \$11,988/YR, ROUGHLY 40% OF ONE TERMINAL SEAT		

All three tiers land inside an existing band rather than inventing one. Desk undercuts every institutional alternative by a wide margin while delivering something none of them carry — a Japanese-language read on the one actor who moves this price deliberately.

The number that matters to the buyer

A subscription is only cheap relative to the capital it is deployed against.

Watch · \$468 a year

On a **\$250,000** trading account that is **0.19% a year**.

One avoided half-per-cent drawdown is \$1,250. **It pays for itself two and a half times on a single mistake not made.**

Signal · \$1,788 a year

On the same account that is **0.72% a year** — less than a typical platform's annual financing cost.

One trade capturing a 1% account move returns \$2,500. **A single correct call per year covers it with change.**

Desk · \$11,988 a year

On a **\$100 million** book that is **1.2 basis points**.

One avoided ten-basis-point loss is \$100,000. **Eight times the annual fee, on a single event, once.**

Where it breaks even, and what holds it there

BREAK-EVEN

318

To cover \$390k before variable costs. **374 at an 85% contribution margin.**

CONSERVATIVE CASE

\$552k

Half the engagement **and** half the conversion — 450 subscribers. Covers the cost base 1.4x.

REPLACEMENT NEEDED

80/mo

To hold 1,800 flat — **958 a year, 0.89% of the engaged pool per month.**

ACQUISITION

near nil

No paid acquisition — onboarding, support, payments and retention still cost.

HOLDING THE BASE FLAT — THE REAL OPERATING CONSTRAINT				
MONTHLY CHURN		REPLACEMENT PER MONTH		
Watch	5.0%	1,186 × 5.0%	=	59.3
Signal	3.5%	571 × 3.5%	=	20.0
Desk	1.0%	54 × 1.0%	=	0.5
79.8 ÷ 9,000 = 0.89%/mo		TOTAL	79.8/mo	· 958/yr
Price is not the constraint. Replacement is. A constant base means replacing cancellations <i>every month, including among the replacements</i> — monthly churn on the standing base, not first-year attrition spread over twelve.				

Publishing does not decay this signal

RELATIVE SIZE

Global FX turnover, daily	\$9,600,000,000,000
USD/JPY turnover, daily	~\$30,000,000,000
1,800 subs at \$50k each	\$90,000,000
share of one day's USD/JPY	0.3%

In a small-cap, publishing your entry moves it before you finish typing. **Publishability is a function of market depth, and this is the deepest market that exists.**

The sequencing rule that protects the franchise: publish first, trade after, same stated delay every time, with the house position shown on every signal including the ones not taken. Boring next to the screenshot accounts — which is the point.

The other business this could be

A licensed research feed does not have to be sold one subscription at a time. That route is real, and it is deliberately not the one this plan starts on.

ROUTE	WHAT IS SOLD	ILLUSTRATIVE REVENUE	WHAT IT REQUIRES FIRST
PRIMARY — THE 90/10 SUBSCRIPTION PATH, COSTED IN FULL			
Retail tiers	Tiered access to the alert and the published record, sold to an audience that already exists	\$2.4M	A track record and a distribution channel. Both are already in hand.
SECONDARY — THE LICENSED ENTERPRISE PATH, ON THE SAME EVIDENCE BASE			
Prototype	5 enterprise contracts, 20 seats, no API	\$0.49M	Paid design partners rather than an interest survey
Institutional	12 contracts, 60 seats, \$0.50M API and data	\$3.08M	Redistribution rights, model validation, insurance, audit
Scaled	25 contracts, 150 seats, \$2.00M API and data	\$11.85M	Everything above, plus regulatory capital and concentration limits

The enterprise numbers are larger, and the enterprise requirements are why they come second. Selling a feed to a treasury means holding redistribution rights, passing model validation, carrying professional indemnity, surviving an audit and answering a client-classification question in every jurisdiction a buyer sits in. Those are Tier 1 and Tier 2 costs, and none are satisfied by having been right. **The subscription path needs a track record and a channel; the enterprise path needs a legal and operational apparatus.**

DECISION — The 90/10 subscription model is the plan of record and the basis for every cost and break-even figure in this deck. The enterprise path is carried as a named later-stage option, gated on the same evidence, and is not costed or priced into anything here.

Ten working days

This is not a platform build. It is a data pipeline, a scoring function and a phone alert.

TO GATE 1

10 days

Principal-directed, agent-written

MVP BUILD COST

\$2–6k

Platform, tokens, domain. No salaries.

PAID MARKET DATA

\$0

Twenty free official sources

RUNNING COST

\$180–400

Per month until subscribers exist

The honest scope. Twenty public endpoints, one table with two timestamps, five weighted sub-scores, a Japanese-language classifier, and a Telegram message. Every one is a solved problem. What is *not* solved is whether the score leads the market — and that costs ten days to answer, not two quarters.

What the platform layer already hands you

Each of these would otherwise be a week of infrastructure work.

Entities & store

Define the collections and get the database, schema and CRUD layer at once. The bitemporal discipline is two fields, not an architecture.

Scheduled functions

Server-side jobs on a cron. Every collector is one function with a schedule.

Outbound HTTP

Backend functions call any endpoint. FRED, ECB, CFTC, Sakana, Telegram and WhatsApp are all the same shape.

Model calls

LLM invocation with structured-output schemas, so the classifier returns a typed object rather than prose to be parsed.

Auth & tiers

User accounts, roles and gated pages — which is the entire subscriber portal minus the copy.

Hosting & domain

Deployed on save, custom domain attached. No environment to stand up, nothing to keep patched.

Three things no platform decides for you: which sources to trust, how the five sub-scores are weighted, and what threshold is worth waking someone for. Those are judgement, and they are where the ten days actually go.

The ten days, gate by gate

01	Day 1	Entities and the seed Four entities, two timestamps on every row — <code>observed_at</code> and <code>known_at</code> . Load the 140 weekly rates and 28 events as the seed.
02	Day 1	The leakage rule One helper every query routes through. Write a deliberately cheating query and confirm it comes back empty before loading anything else.
03	Day 2–3	Twenty collectors One scheduled function per source, each idempotent, each stamping its own fetch time.
04	Day 3	Backfill to January 2022 Six intervention episodes plus the August 2024 unwind. Runs unattended.
05	Day 4	The Japanese reader Sakana AI reads every BOJ and MOF statement in the original and returns a typed object. Lands in a derived table, never the canonical log.
06	Day 5	GATE 0 — score the reader A bilingual reader grades fifty historical statements blind. A day of work that tests the core premise before anything is built on top.
07	Day 6	The scorer and the replay Five sub-indices, fixed weights, plain arithmetic. No model anywhere near it. Then run it over every day since 2022, point-in-time.
08	Day 7	GATE 1 — did it lead? Did the score rise ahead of at least four of six intervention episodes, plus the Aug 2024 unwind?
09	Day 8	Thresholds and the console Alert levels calibrated with a hard budget — no more than two tier-1 alerts a month. One page, phone first.
10	Day 9	The notifier Telegram bot with inline APPROVE / DECLINE / HALVE , WhatsApp for broadcast. Five lines per alert, no exceptions.

What the ten days cost, and what comes after

LINE	DETAIL	MVP
Application platform	Build tier for the ten days, then a running tier	\$50–200/mo
Sakana AI	Japanese classification. Four years of documents cost about \$6 a pass	\$25–80
Analytical model	Case, challenge, copy. Backfill plus ten days live	\$80–250
Market data	Twenty official sources, all free	\$0
Telegram / WhatsApp	Bot API free; Cloud API free service conversations then per-conversation	\$0–40
Domain & mail	One domain, transactional sending	\$40–80
Bilingual grader	Gate 0 — one day, fifty statements, blind	\$300–600
Principal-directed	Ten days of your own direction, no engineer	\$2–6k

PHASE 1	Live signal Trigger: gate 1 passed on four of six. Tick data starts here — a commercial FX API at \$99–299/mo, not a terminal.	\$8–20k · 3 wks
PHASE 2	Subscriber product Trigger: roughly 90 paid from the trial and waiting list.	\$10–25k · 3 wks
PHASE 3	Osaka observer Trigger: three months live with a documented miss list — that gap is the job description.	\$40–90k/yr
PHASE 5	Institutional tier Trigger: 20+ top-tier subscribers and three asking for an API. The point a terminal seat finally pays for itself.	\$25–60k

Rigour is cheap. Scale is expensive.

An independent review priced a staffed desk at \$3.0m a year and a bounded programme at \$250,000. Both are right for what they describe. Neither is the cost of finding out whether this works.

The reconciliation, in one line. The review did not say lean is impossible. It said **lean is not institutional** — and those are different claims. Almost everything it demands on evidence costs nothing but discipline. Almost everything that costs real money buys *capability*, not credibility.

BUSINESS	WHAT IS CLAIMED	WHAT WOULD PROVE IT WRONG	PRICE OF THE ANSWER
A · Trading edge	Original-language interpretation improves option P&L beyond available feeds and simple market variables.	A market-only, English-only or plain phrase-rule baseline matches the full model.	\$15k
B · Risk service	It improves a specific, recurring, expensive decision for a treasury or overlay manager.	Design partners cannot name a decision they would change — or will not pay for the one they can.	\$5k
C · Subscription	An engaged audience converts at a real price and renews past one dramatic episode.	A priced waiting list does not convert, or cohort retention is worse than ~4% monthly churn.	\$2k

One can succeed while another fails. An excellent warning service can be commercially valuable without ever producing a high-return options strategy. Running the three as one business means a failure in any of them looks like a failure in all of them.

Why the order is settled by cost, not preference

BUSINESS

COST TO FIND OUT

COST OF BEING WRONG AT SCALE, PER YEAR

A · Trading edge

staffed desk

200× gap

\$15k

\$3.0m

B · Risk service

product org

60× gap

\$5k

\$300k

C · Subscription

acquisition + service

600× gap

\$2k

\$1.2m

\$1k

\$10k

\$100k

\$1m

logarithmic

Every test costs between sixty and six hundred times less than the thing it decides. That ratio is the whole argument for starting lean — not thrift, but the observation that the answer is nearly free and the commitment is not.

What the reviewers demanded, and what it costs

Every substantive criticism sorted by price. Almost all of it lands in the first block.

REQUIREMENT	COST	WHO DOES IT
FREE — COSTS DISCIPLINE, NOT MONEY. AGENT-EXECUTABLE.		
Pre-register hypothesis, outcome and horizon before testing	\$0	Written once, versioned
Count quiet days, repeated warnings, false alarms, missed events, time under warning	\$0	Agent — arithmetic on the log
Stamp every value with observation date and availability date	\$0	Two columns, decided day one
Six clocks per event; both language feeds captured independently	\$0	An architecture choice
Baselines: phrase-rule, English-only, Japanese-only, market-only	\$0	Four extra runs of existing code
Shadow-log every prediction prospectively from day one	\$0	Agent — the only clean test there is
CHEAP — HUNDREDS TO LOW THOUSANDS, EACH BUYING A SPECIFIC ANSWER		
Intraday executable quotes — turns "the score rose first" into "and a price existed"	\$350–2,400/yr	Entry-tier FX API
Expanded bilingual golden set — negation, quotation, boilerplate, stance change	\$1–3k once	A grader, paid by the day
Counsel scope opinion — one written classification of the actual product	\$2–5k once	One engagement
EXPENSIVE — BUYS CAPABILITY, NOT CREDIBILITY. DEFER UNTIL GATES PASS.		
Staffed team across PM, research, engineering, risk, ops, compliance	\$1.65m/yr	Deferred
Multi-jurisdiction permissions · redistribution data rights · acquisition at scale	\$350k/yr each	Deferred
Trading capital	\$30m+	Deferred

The serious demands mostly cost nothing. They are decisions about how the thing is built, made once, on day one — so a lean programme that makes them

Three tiers, and what each can honestly claim

TIER	ANNUAL	HEADCOUNT	WHAT IT CAN HONESTLY CLAIM
0 · Lean agentic — the research archive	\$2.2–7.7k	none	A reproducible evidence archive, measured baselines, a prospective shadow log, and an honest answer on whether the signal adds anything.
1 · Licensed service — paid design partners	\$47–113k	0.5	A product a treasury can buy, with permissions settled and a fractional human accountable for what ships.
2 · Institutional desk — the reviewer's case	\$3.0m	~8	Returns after implementation cost, at a book where 3% excess return clears the fixed base.

TIER 0 · THE WHOLE OPERATING COST, ITEMISED			
RECURRING		ONE-OFF	
Application platform	\$600 – 2,400	Counsel scope opinion	\$2,000 – 5,000
Sakana · Japanese	\$60 – 150	Bilingual grader	\$1,000 – 3,000
Frontier model	\$600 – 1,500		
Intraday FX quotes	\$350 – 2,400	YEAR ONE, ALL IN	
Domain, mail, monitoring	\$200 – 400	\$5,150 – 15,670	
WhatsApp · Telegram free	\$0 – 480		
Nikkei Digital	\$340		
PER YEAR	\$2,150 – 7,670		
No salaries, no office, no registrations, no trading capital. Agents do capture, classification, comparison, scoring and reporting. One person approves. That is the entire organisation until a gate says otherwise.			

What agents replace — and what they do not

Agents replace this, at Tier 0

Capture — polling twenty sources, stamping arrival, detecting revisions.

Classification — reading the original Japanese, returning a typed object with the triggering phrase.

Comparison — retrieving similar episodes and naming the differences.

Challenge — arguing the opposite case on a separate prompt.

Scoring and reporting — the five sub-indices, the heartbeats, the weekly scorecard.

Reconciliation — checking the record matches what happened.

Agents do not replace this, at any tier

The accountable reviewer. Someone signs. That is the design, not a limitation — a system where no named person approved has no record worth publishing.

The legal classification. Counsel writes it once, about the actual product. An agent cannot hold that opinion or be liable for it.

The bilingual grader. Scoring against a human is the validity test; scoring against another model tests agreement, not truth.

The design-partner conversation. Asking a treasurer what decision they would change is a human conversation, and the answer is the product.

The honest limit of lean. Tier 0 cannot claim returns after implementation cost, cannot sell a licensed product across jurisdictions, and cannot carry a trading book. It is not a smaller institution — it is **the instrument that decides whether the institution deserves to exist.**

The ninety days, re-costed

The reviewer's five gates are right. Their \$250–500k assumed salaried owners, licensed data contracts and multi-jurisdiction legal work.

01	Day 1–15	Define Scope, pre-registered hypothesis and horizon written before anything is tested , evidence schema with both dates on every field, one counsel scope call. ≈\$2-5k
G1	Gate	Legally feasible, and testable Permissible as scoped, and the data needed for a valid test can actually be obtained.
02	Day 16–45	Reconstruct Auditable archive, four baselines, six clocks per event, executable-quote pricing. This is the ten-day build, done properly rather than quickly. ≈\$2-4k
G2	Gate	Reproducible without future knowledge Another person reruns it and gets the same numbers; no availability date is violated anywhere.
03	Day 46–90	Evaluate Frozen holdout. Forward shadow log running from day one, not day forty-six. Expanded golden set graded blind. Four to six priced design-partner conversations. ≈\$3-7k
G3	Gate	Beats the baselines, or does not If the baselines match it, the alpha claim retires and the workflow value is assessed on its own.

LEAN 90 DAYS \$7–16k Same five gates, agent-executed	REVIEWER'S CASE \$250–500k Salaried owner, data contracts, multi-jurisdiction counsel	DIFFERENCE IN RIGOUR none Every evidence requirement met in both	DIFFERENCE IN CLAIM large Lean cannot claim returns, permissions or capacity
---	--	---	---


What unlocks the next tier

MOVE TO	ONLY WHEN	AND NOT BEFORE, BECAUSE
Tier 1 · licensed service	G3 and G4 both pass — the signal beats its baselines and someone has paid for it	Permissions and redistribution rights cost real money and take months; buying them before a customer exists is buying an option on a guess
Tier 2 · institutional desk	Twelve months of prospective shadow log, execution-capacity tests at small size, and a book where 3% excess return clears a \$3m base	At \$30m the hurdle is 10% excess return. The desk is a capital decision wearing a technology costume
Retire the alpha claim	The baselines match the full model	A successful outcome, not a failure. It costs \$15k to learn and saves \$3m a year — and the workflow, archive and risk service may still be worth building

The standard this sets. Be ambitious about the institution and modest about an untested signal. A lean programme that can say "the baselines matched, and we are retiring the claim" is worth more than a funded one that cannot — because only the first was ever capable of telling you something you did not already believe.

Two companies, one product

The Japanese entity produces. The US affiliate distributes. The principal's trading sits outside both.



What each side may publish

The line that makes this work is content, not geography.

RYŪZAN BI KK — YES	RYŪZAN BI KK — NEVER
"Vice-minister used rung 6 wording at 09:47 JST — あらゆる選択肢を排除しない. First use since 11 July."	"Buy three-month 25-delta yen calls."
"Composite stress score 71 → 84. RHETORIC +12, PRICING unchanged."	"Size at 0.6% of book."
"1-month implied vol 11.2, 25-delta risk reversal −1.9 — protection remains cheap relative to the 2024 episodes."	"Invalidation above 158.40."
"Median lead over English-language wires on this event class: 7.0 hours to the London open."	"Enter here. Exit there."
Translated official commentary with the original alongside it.	Any tiering, pricing or promotion of trade outcomes.

One word, struck from every Ryūzan document. In this industry "signals" means trade recommendations, and that is how it reads on a website, in articles of incorporation, or in a licence agreement — whatever was meant by it. The product is a **monitoring feed**, an **official-language classification index**, and **event data**. Free to fix now; expensive to unwind later.

Cost — track A, the entities

Administrative and professional fees only. Converted at ¥154 per USD.

ITEM	JPY	USD	NOTE
Company setup — GK	¥100–300k	\$650–1,950	Lower-cost entity option
Company setup — KK	¥250–600k	\$1,625–3,900	Stronger credibility with banks, clients and immigration
Business Manager specialist	¥200–600k	\$1,300–3,900	Main certificate package for the US partner
Dependent filings	¥100–400k	\$650–2,600	Assumes marriage completed first and documents clean
Founder / shareholder agreement	¥200–800k	\$1,300–5,200	Sale rights, authority, deadlock — essential on a 50/50
Tax and accounting setup	¥50–200k	\$325–1,300	Initial setup only
Translation and document prep	¥50–200k	\$325–1,300	Marriage, birth, banking, capital, support documents
Seals, certificates, misc admin	¥20–80k	\$130–520	Small formation costs
ONGOING			
Monthly accountant	¥20–80k/mo	\$130–520/mo	≈ \$1,560–6,240 a year
TOTALS BY PATH			
Lean GK path	¥720k–1.8m	\$4,675–11,700	Basic company setup and visa preparation
More credible KK path	¥1.07–2.88m	\$6,950–18,700	Recommended where credibility matters
Recommended admin budget	¥1.2–3.0m	\$7,800–19,500	Before capital and operations

Not included, budgeted separately: capitalization (bespoke, set with counsel), office deposit and rent, payroll, relocation, runway, personal tax.

Cost — track B, the technology

Each stage unlocks on evidence. Only the first is needed to answer whether any of this works.

STAGE 0	MVP — the answer machine Twenty free sources, bitemporal store, Japanese classifier, five-part score, four-year replay, console, phone alert, priced waiting list. Also produces the business-plan exhibit and service materials track A's filing needs.	\$2–6k · 10 days
STAGE 1	Live signal Trigger: the replay led on four of six episodes. Tick data starts here — a commercial FX API at \$99–299/mo, not a terminal.	\$8–20k · 3 wks
STAGE 2	Licensable product Trigger: roughly 90 paid. Versioned schema, access control, the generated public record, and the tiering ØDN sells against.	\$10–25k · 3 wks
STAGE 3	Osaka analyst Trigger: three months live with a documented miss list. If this role is also to satisfy the full-time employment requirement it must be scoped full-time and hired earlier — the jobs overlap but are not automatically the same job.	\$40–90k/yr
STAGE 4	Widen coverage Trigger: subscribers asking about the quiet stretches. Extends into Southeast Asia — which is what track A's business plan says the company does.	\$8–20k
STAGE 5	Institutional tier Trigger: 20+ top-tier subscribers and three asking for an API.	\$25–60k

What it costs to be standing, with something that works

ENTITIES · TRACK A

\$7.8–19.5k

KK, filings, founder agreement, tax setup, translation

TECHNOLOGY · STAGE 0

\$2–6k

Ten working days to a measured answer

COMBINED TO STANDING

\$9.8–25.5k

A registered company **and** a working product

CAPITAL REQUIREMENT

\$194.8k

¥30m stated capital under the Oct 2025 rules, plus office, payroll and runway

Read the first three tiles together with the fourth. Under twenty-six thousand dollars buys a registered company, a filed-ready plan and a replayed four-year history — but the ¥30m capital requirement is roughly eight times that, and it is a threshold rather than a fee. **The administrative budget is the small number; eligibility is the large one.**

COST	SITS WITH	WHY IT MATTERS
Technology stages 0-2	Ryūzan BI KK	Deliberate. The product is the KK's asset, so building it creates exactly the operating substance the filing is assessed on
Osaka analyst, office, accountant	Ryūzan BI KK	Payroll and premises are substance, not overhead
Distribution, subscriber marketing, payments	ØDN Ventures	Customer relationships and monetisation sit entirely on the US side
Trading capital, brokerage, premium paid	Personal — neither entity	No contractual link and no performance-linked fee, or the structure reads as asset management

The order to do it in

Track A and track B are not sequential. Running B first is what makes A credible.

00	Week 0	Settle the scope with counsel The content cut, the two-entity structure, the licence direction. This gates the company name, the articles, and every word of the business plan — so it happens before incorporation, not after.
01	Week 1–2	Build stage 0 · in parallel Ten working days. At the end you hold a replay result, a console, a public record and service materials. These are the filing exhibits.
02	Week 1–2	Engage the professionals Immigration specialist, judicial scrivener, accountant, attorney for founder documents. Confirm the capitalization figure against the finished plan.
03	Week 2–6	Incorporate Ryūzan BI KK Articles, seals, registration, tax notifications. The Japan-side partner leads bank communication.
04	Week 3–10	Bank account and funding record Bank timing varies more than anything else here. Keep source of funds and transfers clean and separate from personal accounts from day one.
05	Week 4–10	Operating substance Office, website, service materials, accountant engaged, three-year forecast, hiring plan. Stage 0 supplies most of this rather than describing it.
06	Before filing	Complete the marriage A hard gate on the family path, and the only item that cannot be parallelised or accelerated with money.
06b	Before filing	Satisfy the Oct 2025 eligibility rules ¥30m stated capital, a qualifying full-time employee, Japanese-language capability, applicant experience, and professional verification of the plan. Thresholds, not fees — and the full-time hire cannot wait for month three if the filing depends on it.
07	Week 8–14	Certificate of Eligibility package Business Manager package for the US partner, and dependent applications with marriage, birth, translation and support evidence.

The first five things

01 · Book the scope call

Japanese counsel, one hour, one question: does the content cut keep Ryūzan outside regulated financial services? **Nothing else should be commissioned until that is answered.**

02 · Build stage 0

Ten days, \$2–6k. It costs less than the founder agreement and it determines whether any of the rest is worth doing.

03 · Name the vehicles

Ryūzan for the producer, ØDN Ventures for the distributor. One name per function, and the word "signals" in neither set of documents.

04 · Line up third-party clients

Two corporate consulting engagements — yen-exposure monitoring, or Japan market-entry work — so the KK is not a single-affiliate business on paper.

05 · Fix the capital figure

Bespoke, and the one number deliberately absent from this deck. It comes from counsel once the plan and staffing are final, and it is the largest cash line in the exercise.

Twenty free sources carry the whole MVP

SOURCE	GIVES	CADENCE
PRICE AND POSITIONING		
ECB Data Portal — no key	Daily reference rates to 1999. USD/JPY derived from EUR/JPY ÷ EUR/USD.	Daily
FRED — free key	DEXJPUS, the full Treasury curve, US macro	Daily
Stooq · CFTC COT · FFAJ	Redundant price, speculative positioning, Japanese retail positioning	Daily / weekly / monthly
OFFICIAL SECTOR		
MOF intervention file	Every yen operation with daily detail back to April 1991. What the score is tested against.	Monthly + quarterly
MOF press · BOJ · Federal Reserve	Escalation wording in the original; decisions, vote splits, guidance; FOMC and FIMA repo	Ad hoc / per meeting
TIC · BIS · e-Stat	Japan's Treasury holdings, cross-border yen claims, Japanese CPI and wages	Monthly / quarterly
MARKET SIGNALS — WHAT THE MARKET IS PRICING		
MOF JGB curve · CME FX options	The full Japanese curve; implied vol and 25-delta risk reversals — the price of the convexity you are buying	Daily
Fed funds futures · FRED credit · JPX	What is already priced; HY spreads and equity vol; Nikkei futures and TOPIX Banks	Daily
NEWS, FREE TIER		
GDELT 2.0 · Publisher RSS	Coverage volume including Japanese-language; headline and timestamp from NHK, Jiji, Kyodo, Reuters Japan	15 min / continuous

The intervention record, the policy wording, the positioning, the rate curve and the balance-sheet measure are all free and all official.
 What a terminal buys is speed and convenience — neither of which is what gate 1 is asking about.

The five sub-indices — where each feed lands

SUB-INDEX	BUILT FROM	ANSWERS	WEIGHT
RHETORIC	Sakana output — rung, rate of change, repetition before moving	How close is the official sector to acting?	30%
DIFFERENTIAL	Policy rates, JGB curve, fed funds futures, OIS	Is the reason for the trade widening or closing — and was it priced?	20%
CROWDING	CFTC COT, FFAJ retail, BIS cross-border claims	How many are on the same side, and how fast can they leave?	20%
PRICING	Implied vol, 25-delta risk reversals, HY spreads, TOPIX Banks	What is protection costing, and is stress visible in credit?	20%
FLOW	MOF operations, TIC holdings, FIMA repo usage	Is official money moving, or only official language?	10%

Why PRICING is separate from RHETORIC. They disagree more often than you would expect, and the disagreement is the trade. Officials escalating while risk reversals stay cheap means the market has not believed them yet — precisely when convexity is worth owning. A single blended score would hide it.

Who Sakana AI is

The choice of a Japanese model over a larger American one is the most questionable-looking decision here, and the easiest to defend.

A Tokyo research company

Founded **July 2023**, headquartered in Azabudai Hills, Minato-ku. Builds frontier models shaped to Japanese language and institutions rather than translated into them. Its own framing is "sovereign AI".

The people

Llion Jones, co-founder and CTO, co-authored *Attention Is All You Need* — **the 2017 paper that introduced the transformer**, the architecture every large language model is built on. **David Ha**, CEO, previously led Google Brain's Japan research team.

Who backs it

\$135m Series B, November 2025, at a \$2.65bn valuation — MUFG, Khosla, NEA, Lux, Macquarie, In-Q-Tel. Earlier register: **MUFG, Sumitomo Mitsui, Mizuho, Nomura, Nippon Life, Dai-ichi Life, Meiji Yasuda, Tokio Marine**, plus NVIDIA, Google and Citi.

Read that investor list again, because it is the argument. The institutions that must read Japanese official language correctly every day put their own capital in. **MUFG and Daiwa are named partners, not only investors.** That is a stronger signal than any benchmark score.

What it costs to read four years of Japanese

NAMAZU · PUBLISHED RATES, \$0.95 IN / \$4.00 PER MILLION OUTPUT

BACKFILL — every BOJ and MOF document since Jan 2022

Statements & releases $900 \times 4,000 \text{ tok} = 3.60\text{M}$

Conference transcripts $60 \times 25,000 = 1.50\text{M}$

Input $5.10\text{M} \times \$0.95 = \4.85

Output $960 \times 250 \text{ tok} = 0.24\text{M} \times \$4.00 = \$0.96$

One complete pass over four years = **\$5.81**

Five ladder revisions, re-run each time = **\$29.05**

LIVE $\sim 40 \text{ documents/day} \times 4,000 \text{ tok} = 160\text{k/day} = \$55/\text{yr}$

Reading every word the Bank of Japan and the Ministry of Finance have published in four years, in the original, costs about six dollars. Running it live costs about a dollar a week. The expensive part of this system was never the intelligence.

Sakana publicly lists Namazu at \$0.95 per million input tokens and \$4.00 per million output — the rates above are published pricing, not an estimate.

The manual beat — what a person watches, and where

Two hours a day, 09:00–11:00 JST, where the measurement puts the widest window.

WATCH	WHERE	WHEN (JST)	LOOKING FOR
DAILY			
Nikkei morning edition — 日本経済新聞	Paid app	06:00	Which story leads the market page, and whether the currency moved onto page one
Jiji & Kyodo wires — 時事通信・共同通信	Web / RSS	Continuous	Officials' remarks — usually first, and in the original phrasing
NHK bulletins	Broadcast	07:00 · 12:00 · 19:00	Whether the yen crossed from business into general news
EVENT-DRIVEN — DIARISED IN ADVANCE			
MOF vice-minister remarks — 財務官	Press pen	09:00–10:00	The exact phrase. This is the escalation ladder speaking.
BOJ Governor press conference	Live stream	15:30 on decision days	Tone, hesitation, which question gets dodged twice
Diet Finance Committee — 財務金融委員会	Diet TV	Session days	The Governor and Minister answering questions they did not choose

The four-field form. The words said, **verbatim and in the original**; who said them and in what setting; where and at exactly what time; and how the wording differs from last time. **There is no recommendation field.** The observer supplies evidence. The system supplies the view.

What arrives on the phone

TELEGRAM · TIER-1 ALERT

▲ **STRESS** 71 → 84 14 Sep 09:47 JST

CHANGED MOF vice-minister: rung 6.
"あらゆる選択肢を排除しない"
First use since 11 Jul.
English wires: not yet.

CASE Ladder at 6, spot 3.1 above the
Jul trigger. COT -142k, crowded.

COUNTER Rung 6 held 9 days in Jul with
no action.

TRADE 3m 25-delta JPY call, 0.6% of
book, \$180k, invalid above 158.40

STOP Score below 70, or 10 sessions.

[**APPROVE**] [**HALVE**] [**DECLINE**]

TELEGRAM · 08:30 JST / 19:30 MIAMI

◆ **TOKYO PRE-OPEN** 15 Sep 08:30 JST

STRESS 68 (-3 overnight)
SPOT 153.61 range 152.90–154.50
TODAY BOJ decision Friday. No
projections – statement carries it.
LADDER Rung 5, held 4 days.
WATCH 09:00 press pen · 10:30 JGB
OPEN 1 position. Invalid above
158.40 – 4.79 away.

Nothing requires you this morning.

WHATSAPP · SUBSCRIBER BROADCAST, TIER 2

RYŪZAN SIGNAL · TIER 2

Yen – new escalation, rung 6 of 8.

WHAT WE DID Bought 3m yen upside.

INVALIDATION Above 158.40.

RAISED 09:47 JST **RELEASED** 10:17

This is what we did, not advice for you.

The scheduled message set

Miami gets six kinds of message and no others. **Exactly one ever asks for a decision.**

MESSAGE	FIRES	JST	MIAMI	CHANNEL	ASKS?
Tokyo pre-open	Daily, weekdays	08:30	19:30 prev	Telegram	—
Event alert	On threshold breach	any	any	Telegram	YES · 3 buttons
Overnight Japan	Daily, weekdays	20:00	07:00	Telegram	—
US close handoff	Daily, weekdays	05:00	16:00	Telegram	—
Flow confirmation	MOF release days	19:00	06:00	Telegram	log only
Weekly scorecard	Sunday	Mon 07:00	18:00	Telegram + WhatsApp	—
OUTBOUND TO SUBSCRIBERS — NEVER TO THE APPROVER					
Tiered broadcast	On approval, per tier delay	—	—	WhatsApp + Telegram channels	—

One asking message

Only the event alert carries buttons. If a message has no buttons it needs nothing from you — learnable in a week.

Same shape every time

The field labels never move. The approver reads position, not prose, which is what makes a nine-second read possible.

Heartbeats unconditional


Pre-open and overnight go out on quiet days too. **A missing heartbeat pages the desk before a human notices.**

Numbers carry their age

Every figure prints the time it was true of. A stale number that looks live is worse than no number.

The approval loop

Declines are recorded exactly as approvals are.



A published track record containing only the trades someone took is a marketing document. One containing **every signal the system raised and what was done about it** is evidence — and the record is generated from the store, never written by hand.

Plain English — the terms used in this deck

USD/JPY	How many yen one dollar buys. The number going UP means the yen getting WEAKER.
Basis point	One hundredth of one per cent. Central banks move in 25s.
Negatively skewed	Many small gains, occasional very large losses. The opposite of a lottery ticket.
Yen call	A contract paying if the yen strengthens past an agreed level. You can never lose more than the premium.
Invalidation level	The price at which you accept the idea was wrong. Decided before entry, printed on every alert.
25-delta risk reversal	The price gap between betting up and betting down at equal distance. The cleanest fear gauge this market has.
財務官	Vice Minister for International Affairs — the official who speaks for Japan on the currency. His wording <i>is</i> the ladder.
Rate check	Officials phoning banks to ask where the rate is, without buying. A threat that costs nothing.
JGB	Japanese Government Bond. The long end is the real constraint on how far the BOJ can raise rates.
Bitemporal	Storing two dates on every fact: when it was true, and when it could first have been known.

Carry trade	Borrowing where money is cheap and holding where it pays more, keeping the difference.
The differential	The gap between the two countries' rates. The entire reason the trade exists.
Unwind	Everyone in the same position trying to exit at once. Not a price move — a stampede.
Convexity	A payoff where upside is several times downside. \$180k risked to make \$740k.
Implied volatility	How much movement the options market expects. Higher means protection costs more.
MOF · 財務省	Japan's Ministry of Finance. It orders intervention; the central bank executes.
Intervention	The government buying its own currency with reserves to push the price.
The ladder · rung	The eight fixed phrases officials escalate through before acting. Rung 6 precedes money moving.
COT positioning	A weekly US regulatory file showing how many speculators bet each way. Measures crowding.
Backtest · replay	Running the system over past events. Did the score rise <i>before</i> the move, or alongside it?

What the build hands over

Data model

fx_weekly — 140 rows, ECB reference rates.

events — 28 rows, typed and weighted.

facts — every figure with source and as-of date.

scenarios — generated projections.

Two timestamps on every row: **when it was true of, and when it could first have been known.**

Design tokens

Ground #000000 → #1B1D21. Text #FFFFFF → #5A5F68.

A · flow #3B6FF5 → #21D4FD

B · heat #F5484A → #FFA751

C · gain #12B886 → #63E6BE

D · official #B24BF3 → #F5484A

Archivo 400–900 for display; IBM Plex Mono for every numeral.

Priority order if rebuilding

1. Data layer and facts collection — everything reads from it.
2. The chart and backtest — the most valuable component.
3. Tab shell and tokens.
4. Primer and static views.
5. Projection module — ship without it if no model endpoint.

The rules that must survive

Gradients apply three ways only: clipped to numerals, a 2px card edge, a 3px callout edge — **never as a background fill.**

A number without a source field does not render. Model output never writes to the canonical log. The projected path is never drawn in the same style as the measured series.



Two numbers and a **decision.**

Ten working days produce a measured answer to one question: does the score rise **ahead of** the move, or alongside it? Everything in this deck is downstream of that.

10 days

to gate 1

\$2–6k

to find out

318

subscribers to break even